MSIN0097 PREDICTIVE ANALYTICS INDIVIDUAL COURSEWORK

Can machine learning models predict S&P 500 direction using US multi-asset price and volume data from stocks, cryptocurrencies, and commodities?

Code Repository URL:

https://github.com/maratsakiri/Predictive-Analytics_individual/tree/main/Predictive-Analytics-Individual-TJPQ0

03 MARCH 2026

CANDIDATE NUMBER: TJPQ0

Word Count: 2000

Table of Contents

1. Problem framing 2
   a. Introduction and motivation 2
   b. Dataset and Data Preprocessing 2
   c. Success metrics 2
   d. Operational & Technical Constraints 3
   e. Assumptions and Limitations 3
2. Exploratory Data Analysis 4
3. Data Preparation 7
   a. Train/ Validation/ Test Split Discipline 7
   b. Feature Scaling 7
   c. Data Validation 7
4. Model Exploration and Shortlisting 8
   a. Baseline Model: Logistic Regression 8
   b. XGBoost 9
   c. Neural Network (Multi-Layer Perceptron) 9
   d. LightGBM with Technical Feature Engineering (Logistic Regression Ensemble) 10
   e. CatBoost Classifier 10
   f. Model Comparison 10
5. Fine-Tuning and Evaluation 12
   a. Hyperparameter-Tuning 12
   b. Cross Validation 12
   c. Test Set Performance 12
   d. Feature Importance 13
6. Final Solution 14
   a. Limitations, Risks & Next Steps 14
   b. Model Summary Card 14

APPENDIX 15
   A. Time-Varying Correlations 15
   B. Boxplots for Outlier Visualisation 16
   C. XGBoost feature importance 17
   D. CatBoost Validation Learning Curve 18
   E. Catboost Feature Importance, SHAP and Error Analysis 18
   F. Agent Decision Register 20

References 24

1. Problem framing

a. Introduction and motivation

Stock market prediction is a challenging problem in quantitative finance, with applications in portfolio management, risk assessment, and algorithmic trading. This project investigates whether multi-asset signals, combining traditional equities, cryptocurrencies, and commodities, can predict next-day S&P 500[1] direction.

The hypothesis is that cross-asset relationships contain predictive information (e.g. Bitcoin's rise might signal risk appetite, Gold's surge could indicate flight to safety) and leveraging machine learning can uncover exploitable patterns in the market structure.

Binary classification will be used to predict the S&P 500's daily direction, labelling instances as 1 (UP) if the next-day close exceeds today's or 0 (DOWN) otherwise. This approach prioritizes actionable trading signals over price magnitude, managing volatility more effectively than regression. The observations contained follow a 54.2%(UP)/45.8%(DOWN) distribution, indicating an acceptable slight imbalance reflecting historical bull market bias.

b. Dataset and Data Preprocessing

The dataset used is the US Stock Market Data from Kaggle, containing 1,242 daily observations of 38 columns (variables). The asset classes involved are stock indices (S&P500, Nasdaq100), individual equities (Apple, Tesla, Microsoft, Google, Nvidia, Berkshire Hathaway, Netflix, Amazon, Meta), cryptocurrencies (Bitcoin, Ethereum) and commodities (Gold, Silver, Platinum, Crude Oil, Natural Gas, Copper). Each asset includes price (close) and volume features. The dataset spans the five-year period of 2019-2024, encompassing major market events like COVID-19 crash (2020), cryptocurrency boom (2021), inflation surge (2022), and AI (2023).

The data demonstrates high quality, requiring only a brief data cleaning procedure. The dataset exhibits a low overall missing value rate of 2.57%, with complete price records achieved through forward-filling (Codefinity, 2023). Platinum volume was excluded due to a substantial (48.8%) data gap. No duplicate observations were detected, commas were removed from all numeric fields, and a binary target variable was constructed (1=UP if next-day price exceeds today's price, 0=DOWN otherwise).

c. Success metrics

The primary evaluation metric is ROC-AUC[2], targeting scores above 0.55. ROC-AUC measures the model's ability to discriminate between UP and DOWN days across all classification thresholds, where 0.50 represents random chance and 1.0 perfect separation (EvidentlyAITeam, 2024). This threshold-independent metric remains robust to the dataset's class imbalance,

---

[1] S&P 500 is an index that tracks the share prices of 500 of the largest publicly traded companies in the US including leading companies across a wide range of industries such as Amazon, Netflix, Apple, and Microsoft (HSBC, 2025).

[2] Area Under the Receiver Operating Characteristic Curve

unlike accuracy which can be misleading when an "always predict DOWN" strategy would achieve 54% accuracy despite being uninformative. Secondary metrics, including Precision-Recall AUC, F1-Score, and Directional Accuracy, supplement ROC-AUC with threshold-specific diagnostics.

d. Operational & Technical Constraints

- Performance: Inference must occur in <1 second to meet real-time trading requirements.
- Integrity: Models require feature importance for regulatory trust, strict temporal validation to eliminate lookahead bias, and fixed random seeds (SEED=42) for full reproducibility.
- Cost Sensitivity: False positives result in direct financial loss, hence higher precision is preferred over recall.

e. Assumptions and Limitations

The project's foundation rests on weak-form market efficiency, assuming that technical patterns and cross-asset correlations, specifically between Bitcoin, Gold, and equities, hold predictive value for 24-hour horizons. While historical pattern persistence is assumed, the strategy acknowledges volatility during regime shifts and accepts the inherent survivorship bias in using accurate, currently-trading data.

Critical Limitations:

- Scope Restrictions: The 2019-24 timeline includes disruptive events like the COVID-19 crash, and columns are limited to price/volume, excluding macroeconomic, sentiment data.
- Methodological Simplification: Binary classification ignores the magnitude of price movements.
- Operational Realism: Real-world performance is likely lower than back tested results, since spreads, slippage, or transaction fees are not considered in trading simulations.
- Selection Bias: The focus on nine large-cap tech stocks excludes small-cap, value, and international market dynamics.

2. Exploratory Data Analysis

Exploratory data analysis reveals critical patterns and relationships within the multi-asset dataset that inform feature engineering and model selection strategies.

Univariate analysis of individual assets demonstrates substantial variation in price ranges and volatility profiles across asset classes. Stock indices show concentrated, approximately unimodal distributions; S&P 500 clustering between 2,500–5,000 and Nasdaq100 displaying a slight bimodal structure reflecting its 2022 trough and recovery (Figure 1). Technology equities exhibit wider, irregular distributions, with Tesla and Nvidia both showing bimodal patterns driven by stock splits and rapid price appreciation respectively (Figure 2). Distribution analysis confirms non-normality across all variables (Shapiro-Wilk $p<0.001$), with positive skewness in growth stocks and cryptocurrencies.

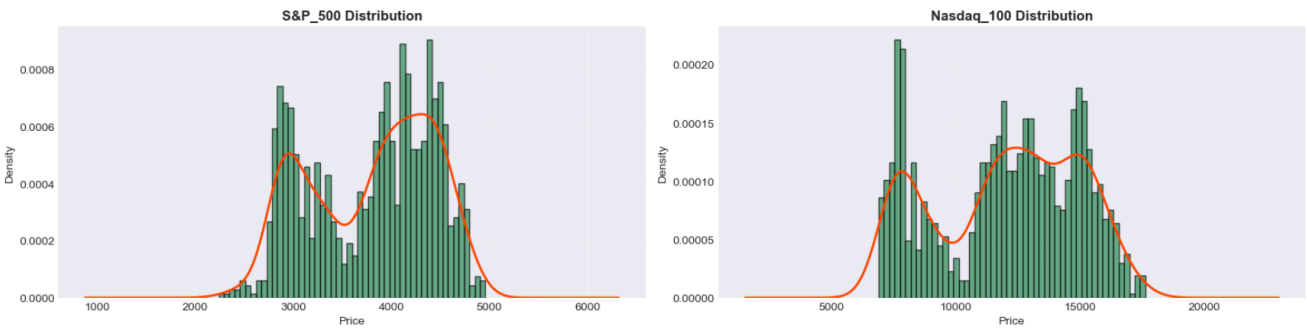

Figure 1: Price Distributions Across Stock Indices

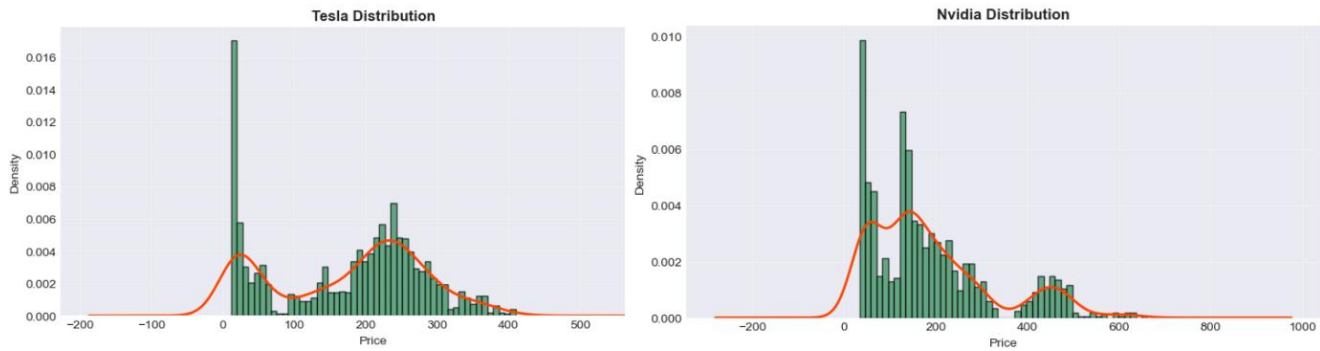

Figure 2: Price Distributions Across Technology Equities

Temporal pattern analysis identifies several regime changes critical for the model validation strategy (Figure 3). The dataset encompasses five distinct market phases: pre-pandemic stability (February 2019-March 2020), COVID crisis (March-June 2020), recovery and cryptocurrency boom (July 2020-November 2021)(Sweeney,2026), inflation surge and bear market (December 2021-2022), and AI-driven recovery (November 2022-February 2024). Bitcoin demonstrates dramatic multi-month drawdowns during the 2022 collapse, while crude oil's brief negative pricing event in April 2020 is visible as an extreme spike. Gold exhibits characteristic safe-haven appreciation during market stress periods (Salisu/Raheem, 2021).

Figure 3: Temporal Price Trends and Market Regime Changes (2019-2024)

Volatility analysis reveals regime-dependent patterns across asset classes (Figure 4). Bitcoin exhibits sustained extreme volatility, frequently exceeding 150% annualised during the 2021-22 period. S&P500 and Gold volatility spikes are largely coincident, both peaking sharply during the COVID crash before reverting, suggesting limited diversification benefit between the two during systemic stress events.

Figure 4: 30-Day Rolling Volatility (Annualized)

Cross-asset correlation analysis (Figure 5) reveals strong multicollinearity among technology stocks (>0.85), with S&P 500 and Nasdaq100 near-perfectly correlated (0.97). Bitcoin shows moderate equity correlation (0.84) but this is highly time-varying, ranging from +0.90 during bull markets to -0.65 during crypto stress (Appendix A). Gold oscillates between -0.75 and +0.80 with S&P500, undermining its safe-haven reliability.

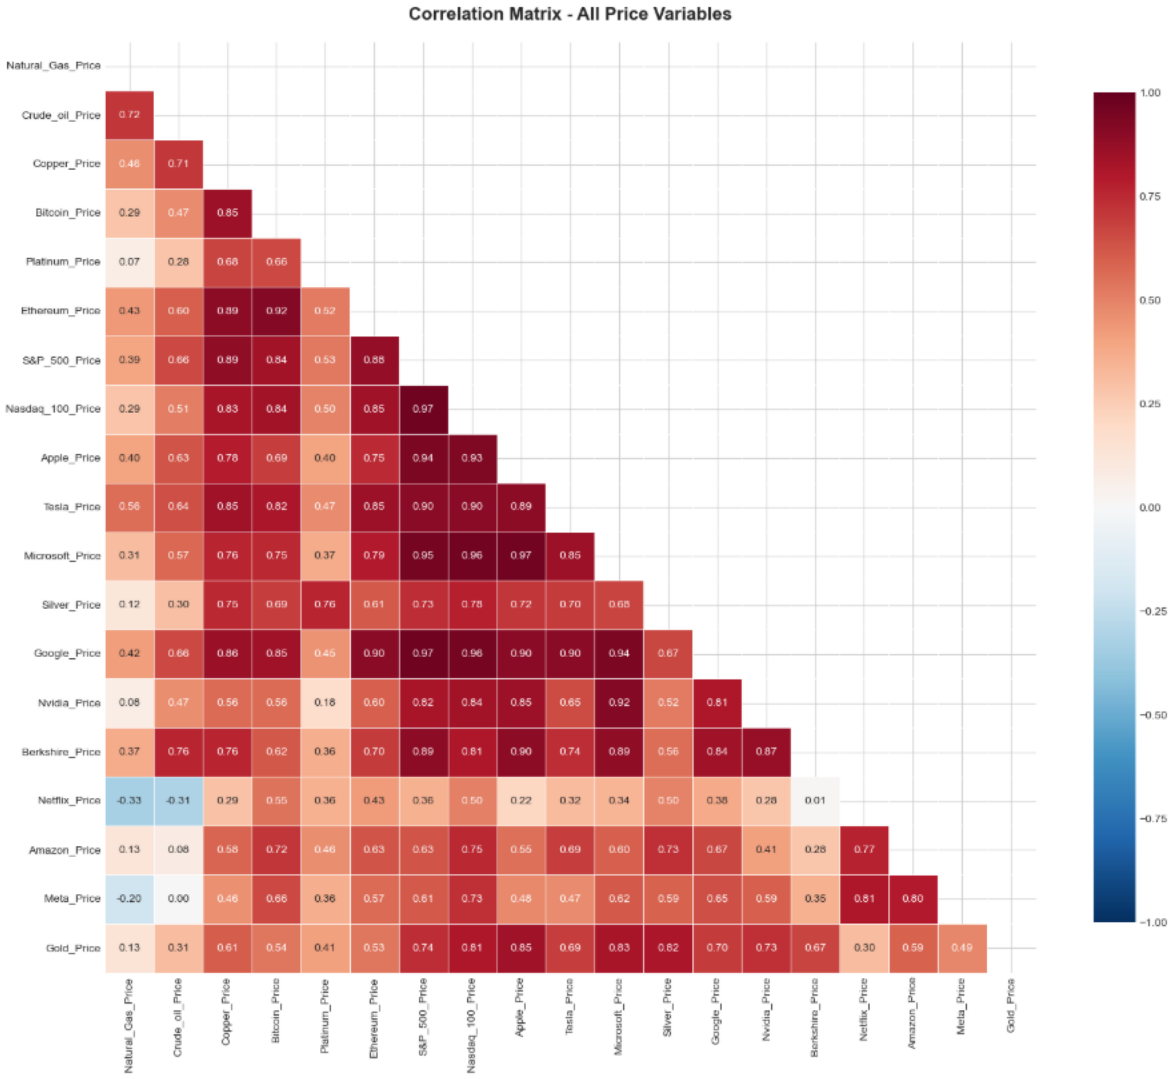

Figure 5: Cross-Asset Correlation Matrix

3. Data Preparation

a. Train/ Validation/ Test Split Discipline

The dataset is partitioned chronologically into training (60%, February 2019-December 2021), validation (20%, January 2022-August 2022), and test (20%, September 2022-January 2024), strictly preserving temporal ordering to prevent lookahead bias. The validation set encompasses the 2022 bear market, providing model selection under stressed conditions, while the test set captures the subsequent recovery.

b. Feature Scaling

Features are standardised to ensure comparable scales across assets with vastly different price ranges (Bitcoin in tens of thousands versus crude oil in tens). The scaler is fitted exclusively on training data, with learned parameters applied to validation and test sets to prevent leakage. The specific scaling method varies by model and is documented in Part 4.

c. Data Validation

Split integrity is verified through assertions confirming no temporal overlap, correct chronological ordering, and consistent feature dimensions. The final dataset comprises 745 training, 248 validation, and 245 test observations. Boxplot analysis reveals outliers in volume features (3-5% of observations), concentrated during high-volatility events and retained as genuinely informative market signals (Appendix B).

4. Model Exploration and Shortlisting

Model selection follows a three-stage progression: linear baseline (Logistic Regression), non-linear architectures (XGBoost and MLP), and regularised gradient boosting ensembles (LightGBM+LR, CatBoost). Each stage is motivated by the evaluation of the previous. All five models are assessed on validation ROC-AUC (target:0.55), overfitting gap, and interpretability. Classical time series methods (ARIMA) and recurrent architectures (LSTM) were not pursued, as multivariate feature engineering with gradient boosting better exploits cross-asset signals than univariate sequential models.

a. Baseline Model: Logistic Regression

Logistic Regression estimates next-day S&P 500 direction as a weighted sum of input features. Its coefficients directly quantify each feature's directional influence, a transparency requirement in regulated financial applications. Ananthakumar and Sarkar (2017) validate logistic regression for stock classification using multi-asset features, supporting its use as a benchmark. With L2 regularisation, the model achieves validation AUC of 0.5348, confirming limited but genuine linear signal across cross-asset features, while falling short of the 0.55 target.

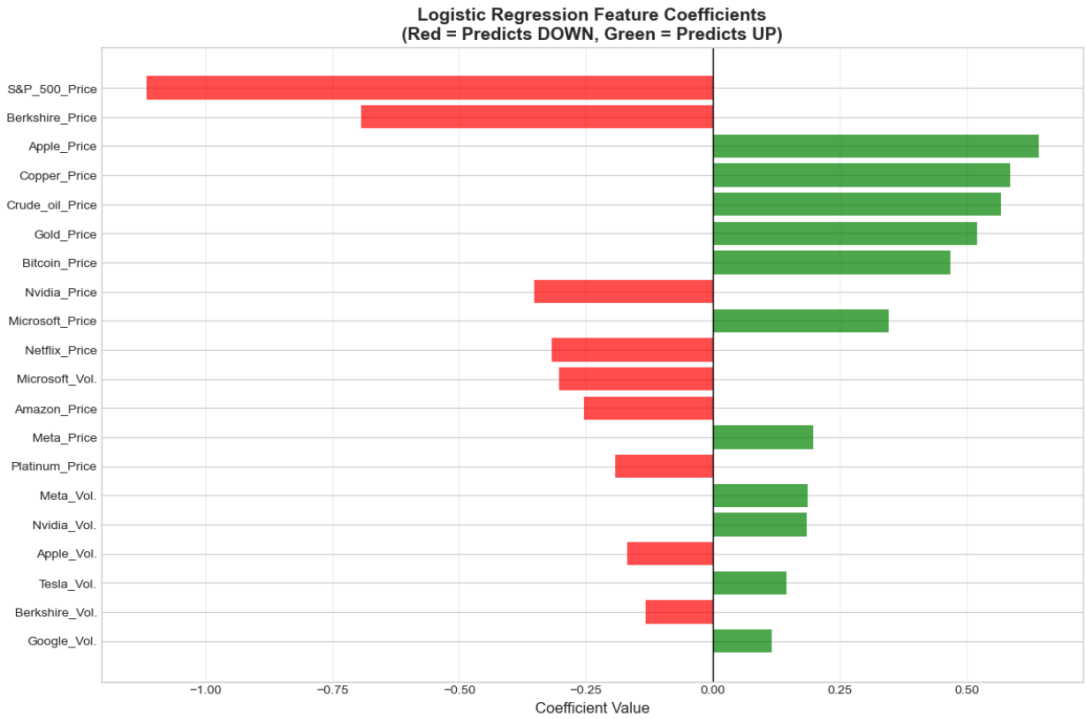

Figure 6: Logistic Regression Baseline Model - Feature Coefficients and Directional Influence

b. XGBoost

XGBoost builds trees sequentially employs gradient boosting to correct residual errors. Nair et al. (2025) validate it for next-day stock direction on tabular financial data. However, on the raw 37-feature set, XGBoost suffers severe overfitting (TrainAUC:1.00, ValAUC:0.4850) even with conservative regularisation (max_depth=3,learning_rate=0.1,subsample=0.8), confirming that raw price and volume features are insufficient for deep tree ensembles and motivating technical feature engineering (Section 4.e). Feature importance (Appendix C) identifies Tesla price, Meta volume, and Google price as the dominant predictors.

c. Neural Network (Multi-Layer Perceptron)

A two-layer MLP (64 and 32 neurons, ReLU) tests whether neural representation learning captures non-linear patterns. Following evidence of competitive MLP performance on engineered financial features (Abolmakarem et al., 2024), the architecture achieves validation AUC of 0.4956 with stable training loss (Figure 7) and minimal overfitting, but only marginally exceeds the linear baseline, confirming that capacity alone without richer features is insufficient.

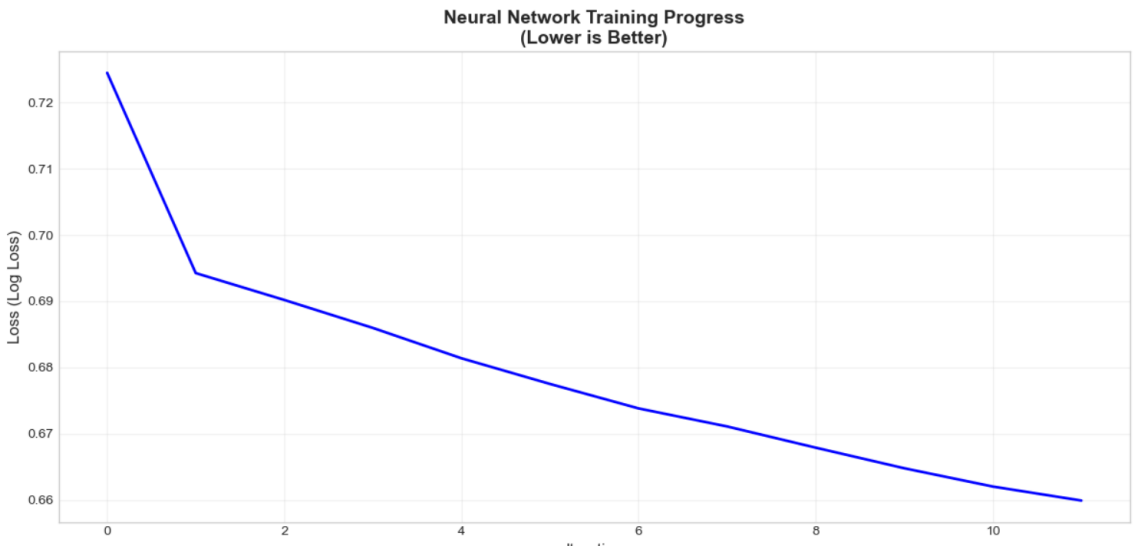

Figure 7: Neural Network Training Progress (Multi-Layer Perceptron)

d. LightGBM with Technical Feature Engineering (Logistic Regression Ensemble)

Twenty-four technical indicators (moving averages, momentum signals, rolling volatility, and cross-asset ratios) extend the feature set from 37 to 60+ variables. Yuvaraj et al. (2025) confirm gradient boosting with domain-driven features as a robust approach for market data. LightGBM is regularised (num_leaves=20, reg_alpha=0.1,reg_lambda=1.0,early stopping), with learning curve convergence shown in Figure 8. Despite this, severe overfitting persists (TrainAUC" 0.9835); a 50/50 LR ensemble partially mitigates this, reaching validation AUC of 0.5322.

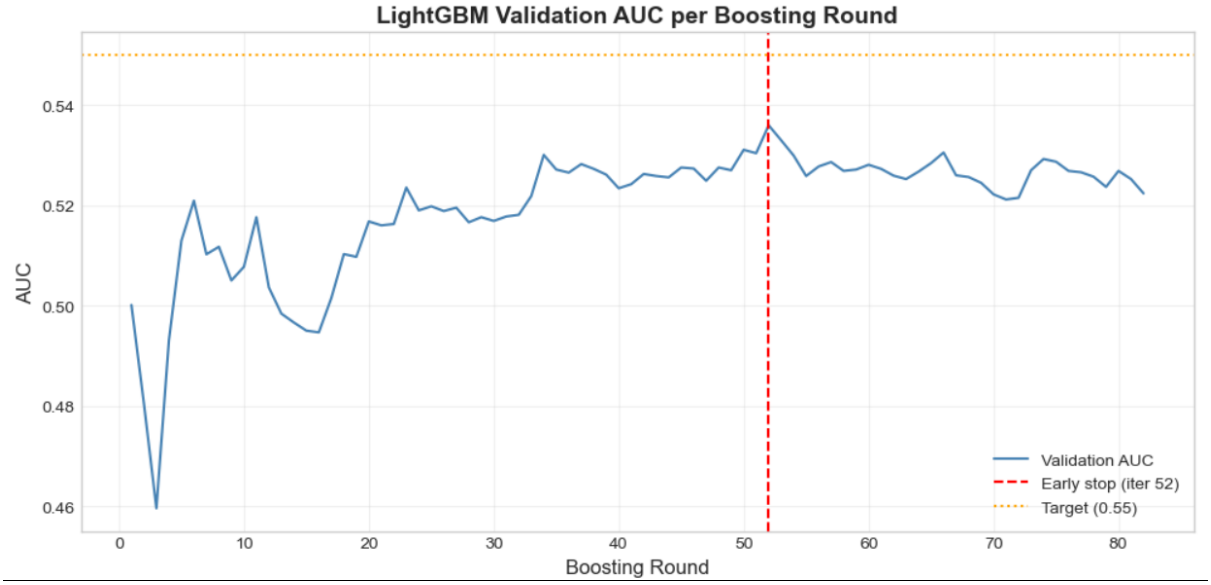

Figure 8: LightGBM Validation AUC per Boosting Round with Early Stopping

e. CatBoost Classifier

CatBoost's ordered boosting algorithm processes training data sequentially to prevent target leakage, while its symmetric tree structure provides more uniform regularisation than LightGBM's leaf-wise growth (GeeksforGeeks, 2024). The same 60+ feature set is used with early stopping (depth=4,l2_leaf_reg=10,learning_rate=0.05). A 50/50 LR ensemble is applied for consistency, with standalone and ensemble results compared in the next section.

f. Model Comparison

Table 1 and Figure 9 summarise validation performance. No model meets the 0.55 target, reflecting the fundamental difficulty of next-day direction prediction under weak-form efficiency. Logistic Regression (0.5348) establishes a competitive linear baseline with a small overfitting gap. XGBoost collapses to 0.4850 due to training memorisation (Train AUC 1.00). The MLP avoids overfitting but adds no meaningful signal over the linear baseline (0.4956). Technical feature engineering substantially narrows the gap: LightGBM+LR reaches 0.5322, though residual overfitting remains. Standalone CatBoost achieves the highest validation AUC (0.5404) with a contained overfitting gap (0.1040); adding LR to CatBoost reduces performance (0.5245), suggesting ensemble averaging introduces unnecessary bias here. CatBoost is selected as the

final model, maximising validation ROC-AUC while maintaining the lowest overfitting gap among the gradient boosting candidates.

Model	Train AUC	Val AUC	Val Acc	Val F1	Overfit gap	Target (0.55)
CatBoost	0.6444	0.5404	0.4672	0.6286	0.1040	NO
Logistic Regression	0.6116	0.5348	0.4597	0.6193	0.0768	NO
LightGBM + LR Ensemble	0.9835	0.5322	0.4590	0.5510	0.4514	NO
CatBoost + LR Ensemble	0.6253	0.5245	0.5164	0.6014	0.1008	NO
MLP Neural Network	0.5187	0.4956	0.4476	0.6184	0.0231	NO
XGBoost	1.0000	0.4850	0.4194	0.5636	0.5150	NO

Table 1: Model Performance Comparison on Temporal Validation Set (20%)

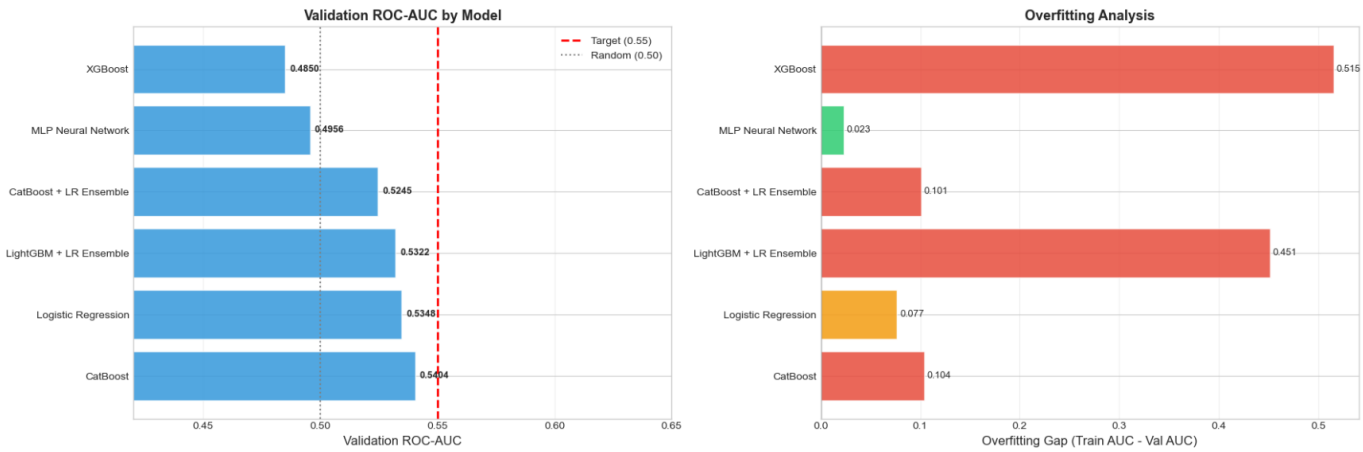

Figure 9: Model Performance Comparison and Overfitting Analysis

5. Fine-Tuning and Evaluation

a. Hyperparameter-Tuning

Hyperparameter optimization employs grid search with 5-fold cross-validation, evaluating 27 XGBoost configurations across learning rate, maximum depth, and estimator count. The optimal configuration (learning_rate=0.1,max_depth=5,n_estimators=150) yields a mean cross-validation ROC-AUC of 0.5389, establishing a tuning benchmark. The selected final model remains CatBoost from Part 3.

b. Cross Validation

5-fold stratified cross-validation on CatBoost yields a mean ROC-AUC of 0.4927 (±0.0275), with fold scores ranging from 0.4571 to 0.5408. High variance across folds reflects the instability of short financial time series, where different temporal windows capture different market regimes.

c. Test Set Performance

Test set evaluation (09/2022–01/2024, n=245) yields ROC-AUC of 0.4826, below both target (0.55) and random baseline (0.50), validating market efficiency. ROC and Precision-Recall curves (Figure 10) confirm performance marginally above chance across most thresholds, with PR-AUC of 0.5371. The confusion matrix (Figure 11) reveals a systematic prediction bias toward UP (153 of 245 predictions), with precision of 52.9% against a base rate of 55.1%, indicating minimal discriminative power beyond the class prior. Error analysis (Appendix E) shows errors clustering in May and September 2023, coinciding with sharp intraday reversals where momentum indicators lag structural shifts.

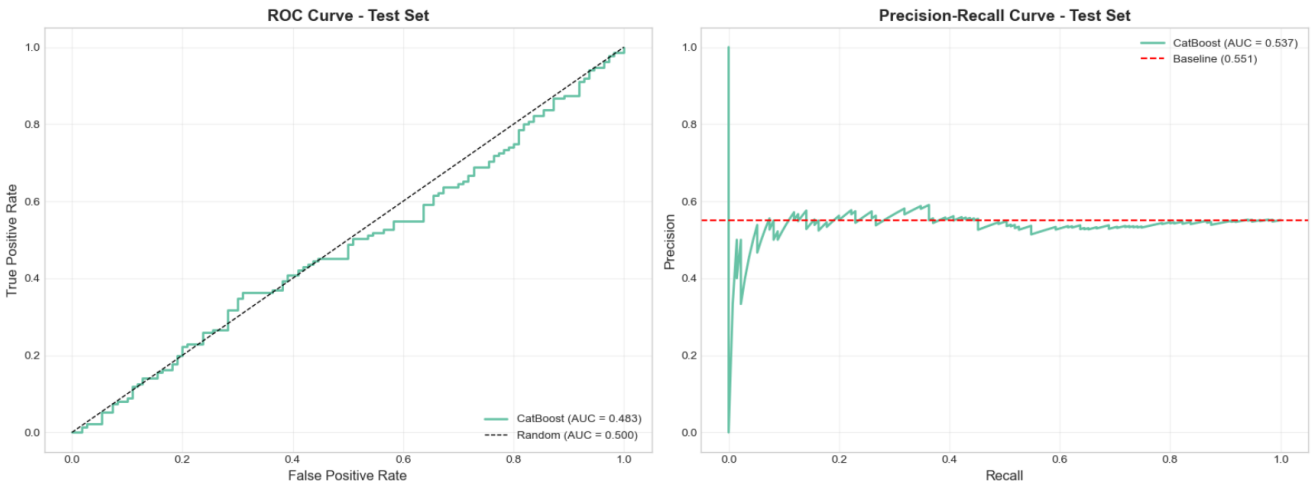

Figure 10: ROC and Precision-Recall Curves - CatBoost Test Set Evaluation

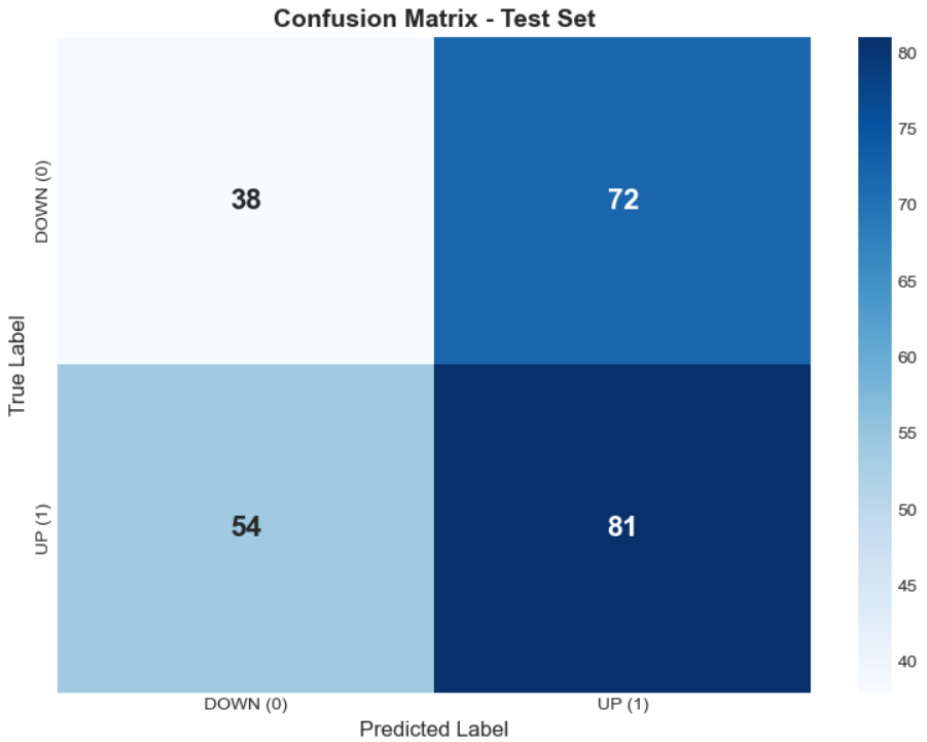

Figure 11: Confusion Matrix - CatBoost Test Set Predictions (n=245)

d. Feature Importance

CatBoost feature importance identifies SP500_Return_10d (22.3%) as the dominant predictor, followed by Nasdaq_100 volume (11.7%) and Amazon volume (9.3%), confirming medium-term momentum as the primary learnable pattern (Appendix E).

The test AUC of 0.4826 represents a decline from validation (0.5404) and falls below the random baseline, consistent with the efficient market hypothesis: price-derived signals are transient and erode out-of-sample. Without fundamental, sentiment, or macroeconomic data, technical indicators on daily horizons capture noise as much as signal.

6. Final Solution

This project investigated whether ML models can predict next-day S&P 500 direction using multi-asset price and volume data. The answer is partially, but not reliably: cross-asset technical features carry weak directional signal that does not persist out-of-sample, with test ROC-AUC declining to 0.4826, consistent with weak-form market efficiency and the limits of price-derived features across regime shifts. CatBoost is selected as the final model, prioritising generalisation stability over inflated in-sample performance. Despite a marginally lower validation AUC than LightGBM+LR, its substantially more contained overfitting gap (0.1040vs0.4513) justifies the choice. SP500_Return_10d dominates feature importance (22.3%), confirming medium-term momentum as the primary learnable signal. The 0.55 ROC-AUC target was not achieved; with basic technical features, market efficiency constrains predictability below the threshold for reliable directional forecasting.

a. Limitations, Risks & Next Steps

The 2019-24 training window introduces a bull-market bias that likely inflates validation performance relative to the bear-market test period. The binary framework ignores move magnitude, and the absence of fundamental and macroeconomic data limits the model's ability to anticipate regime changes. In deployment, concept drift risk is significant; the model requires periodic retraining and should not be used as a standalone trading signal without additional risk controls. Performance could be improved by incorporating macroeconomic indicators (yield curves, VIX, Fed signals), news and social media sentiment, expanding the training window to pre-2008 regimes, and exploring attention-based architectures (LSTMs/Transformers) for longer-range temporal dependencies.

b. Model Summary Card

Model	CatBoost Classifier
Task	Binary classification: next-day S&P 500 direction (UP/DOWN)
Features	60+ variables: 37 market price/volume + 24 engineered technical indicators
Data	US multi-asset daily data, Feb 2019 – Jan 2024 (Kaggle); 1,242 observations; 60% train/ 20% validation/ 20% test temporal split
Intended use	Research and educational demonstration of ML in financial forecasting
Not intended for	Live trading, investment decisions, or deployment without additional risk controls
Validation AUC	0.5404 (best among 5 candidates; does not meet 0.55 target)
Test AUC	0.4826 (below random baseline - performance degrades out-of-sample)
Key caveat	Model is sensitive to regime change; results should not be extrapolated beyond the 2019-2024 training period
Deployment risk	No concept drift monitoring: requires periodic retraining as market dynamics evolve

APPENDIX

A. Time-Varying Correlations

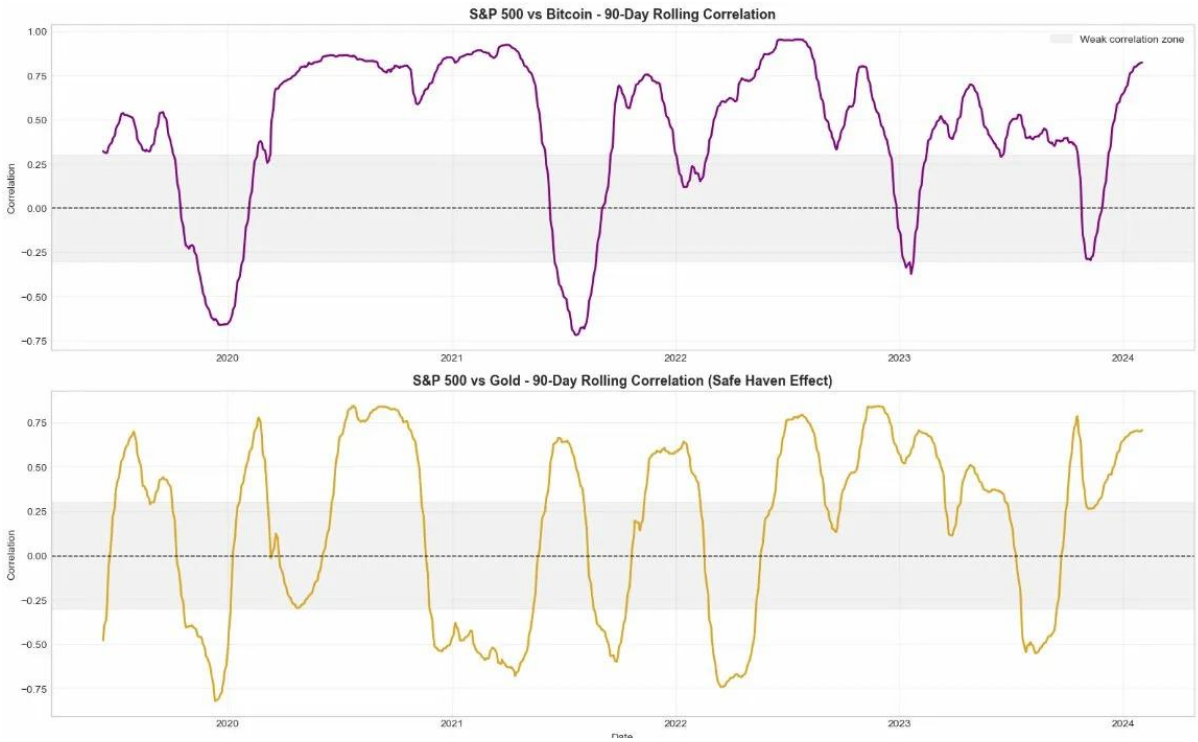

Figure A1: S&P 500 vs Bitcoin and S&P 500 vs Gold 90-Day Rolling Correlation

Rolling correlation analysis reveals substantial time-variation in cross-asset relationships, explaining the moderate static correlations observed in Figure 5. Bitcoin-equity correlation fluctuates dramatically from -0.75 (March 2020 COVID crash, late 2022 crypto winter) to +0.95 (2021 bull market, 2023 recovery), demonstrating risk-on behavior during market rallies and decoupling during stress. Gold-equity correlation similarly varies from -0.75 (traditional safe-haven behavior during 2020 crisis) to +0.85 (simultaneous rally in 2021), invalidating static correlation assumptions. The weak correlation zone (gray band around 0.00) indicates periods where assets move independently. These temporal dynamics underscore the limitations of static correlation matrices for financial forecasting and support the need for adaptive modeling approaches.

B. Boxplots for Outlier Visualisation

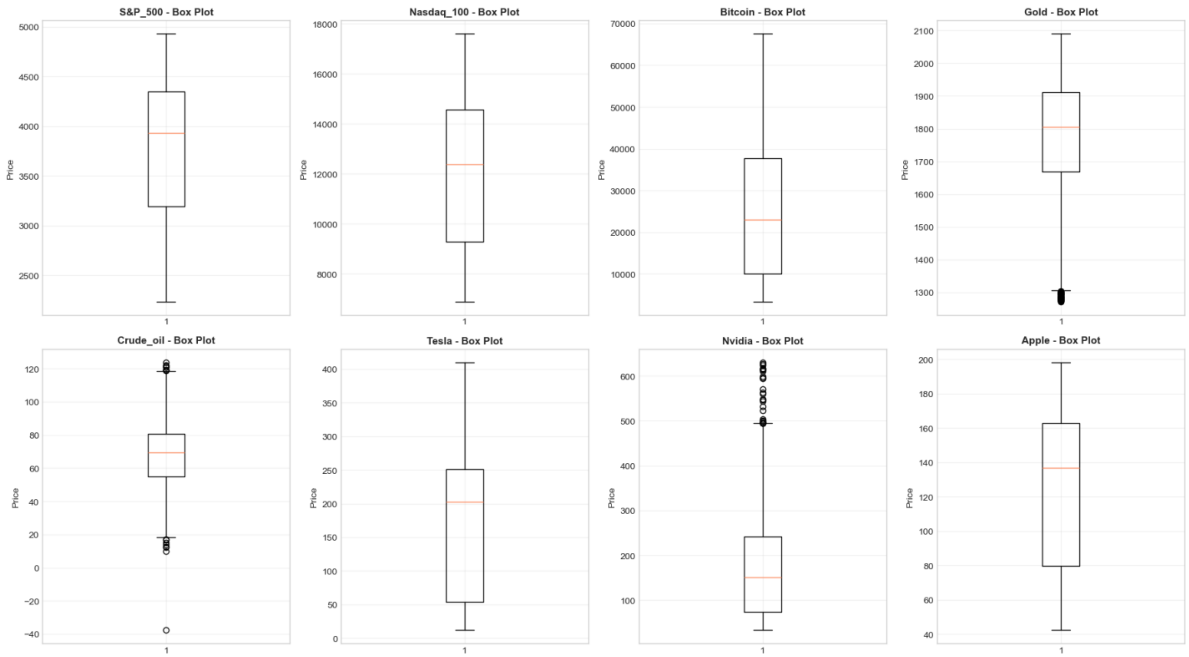

Figure B1: Boxplot Analysis - Outlier Distribution Across Price and Volume Features

C. XGBoost feature importance

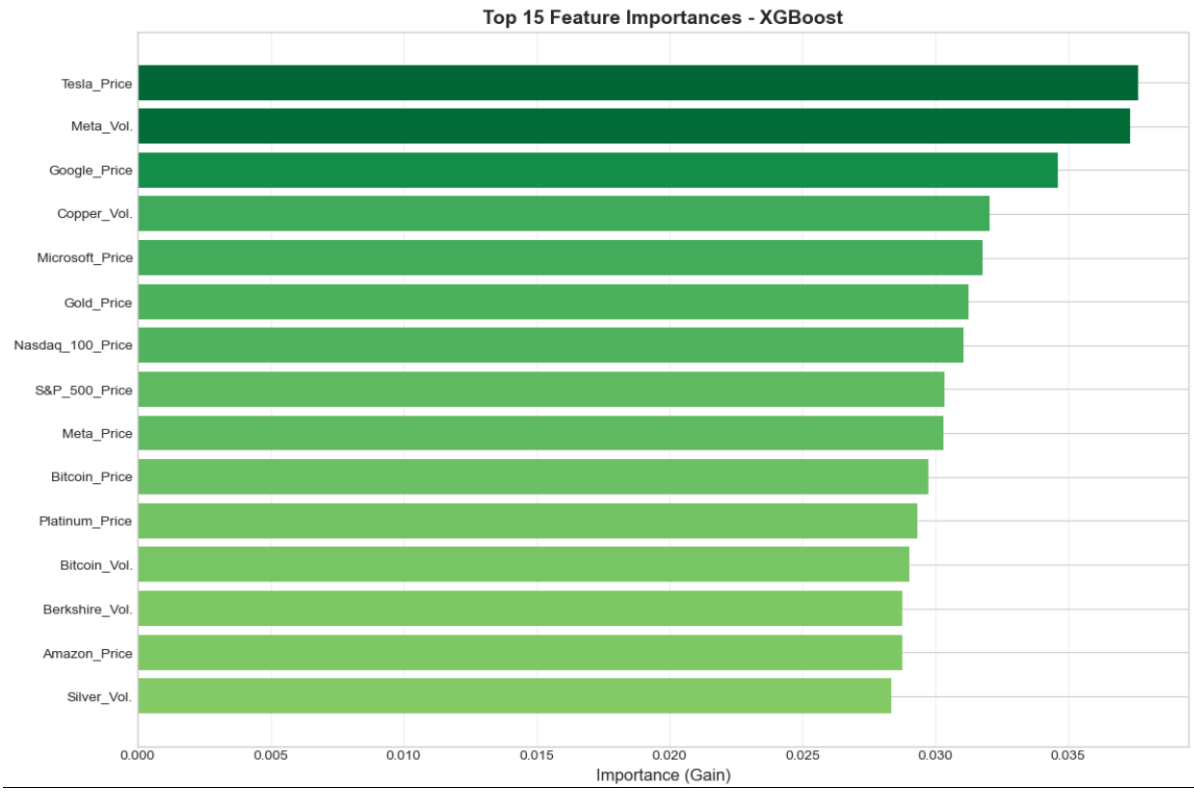

Figure C1: XGBoost Feature Importance (Top 15 Features)

Figure B1 presents the top 15 features ranked by gain importance from the XGBoost classifier. Tesla price (0.038) emerges as the strongest predictor, followed closely by Meta volume (0.037) and Google price (0.035). Copper volume (0.032) and Microsoft price (0.031) complete the top five. Notably, the distribution is relatively balanced across all 15 features, with no single feature exceeding a gain of 0.04. This reflects XGBoost's ensemble nature, which distributes predictive weight across many variables rather than concentrating it in a few. The prominence of individual equity prices and volumes, rather than broad macro indicators, suggests that cross-asset momentum signals carry meaningful information in the model's splits, even if their aggregate contribution to directional prediction remains modest (val AUC: 0.4850).

D. CatBoost Validation Learning Curve

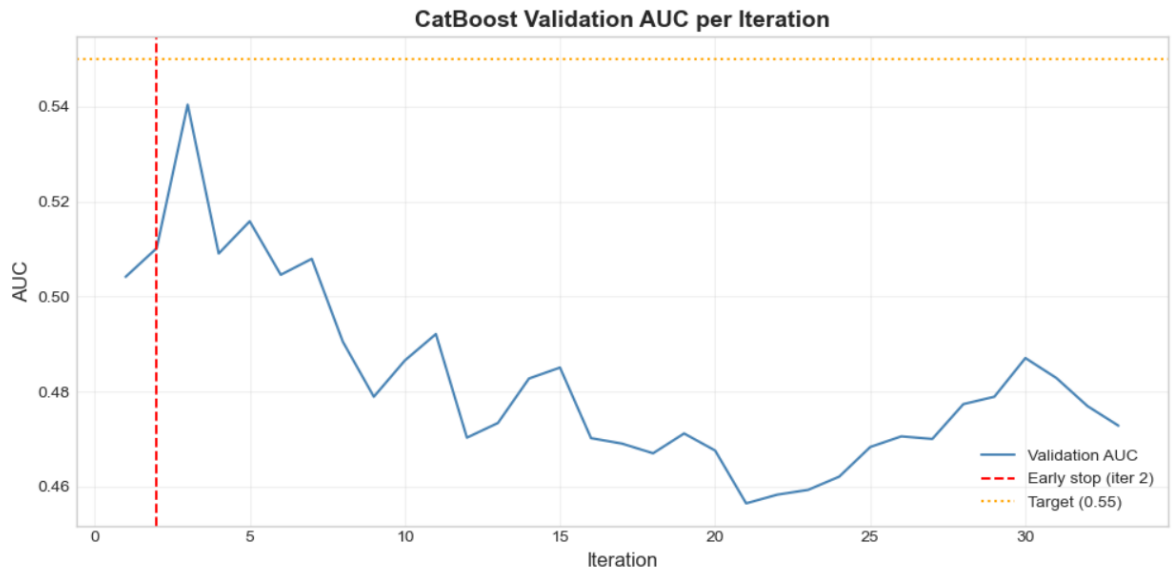

Figure D1: CatBoost Validation AUC per Boosting Round with Early Stopping

E. Catboost Feature Importance, SHAP and Error Analysis

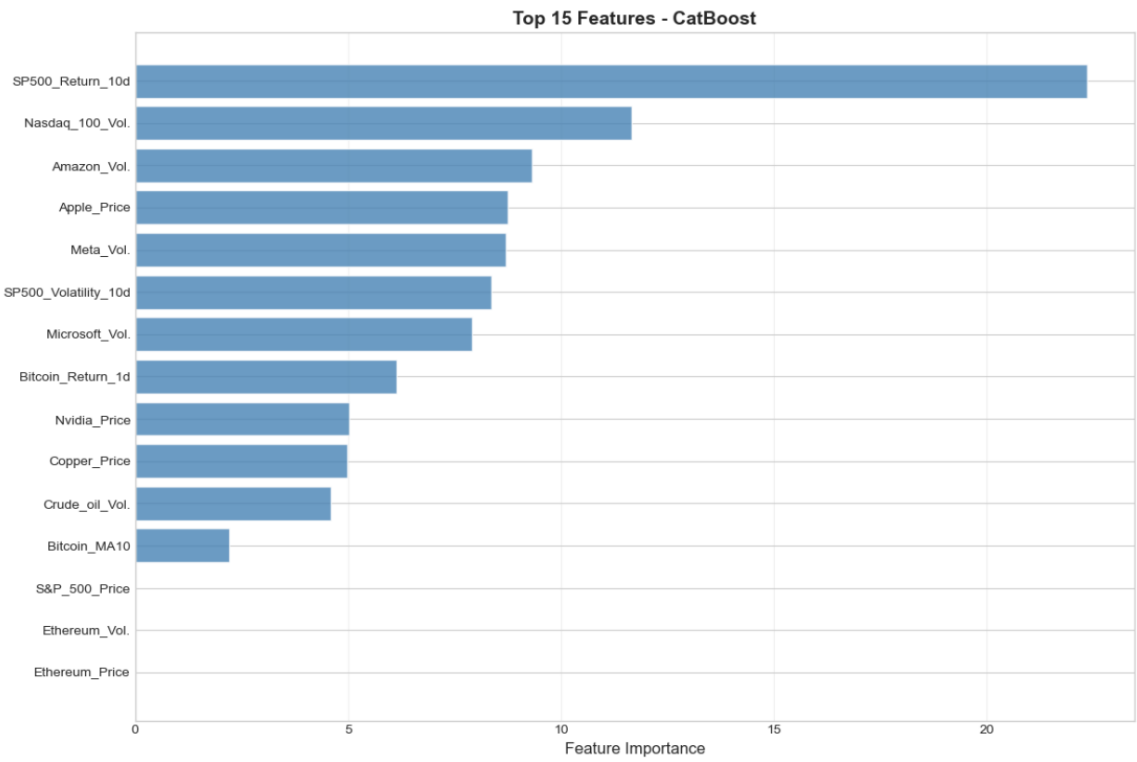

Figure E1: CatBoost Feature Importance - Top 15 Predictors by Split Gain

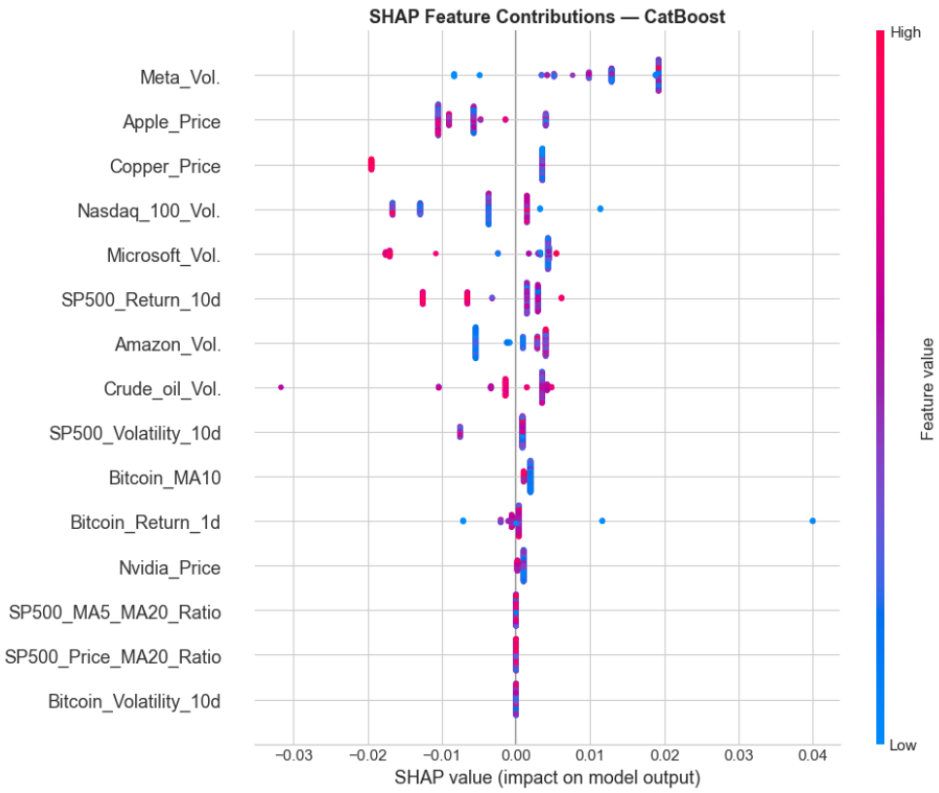

Figure E2: SHAP Summary Plot - Feature Contribution Direction and Magnitude (CatBoost)

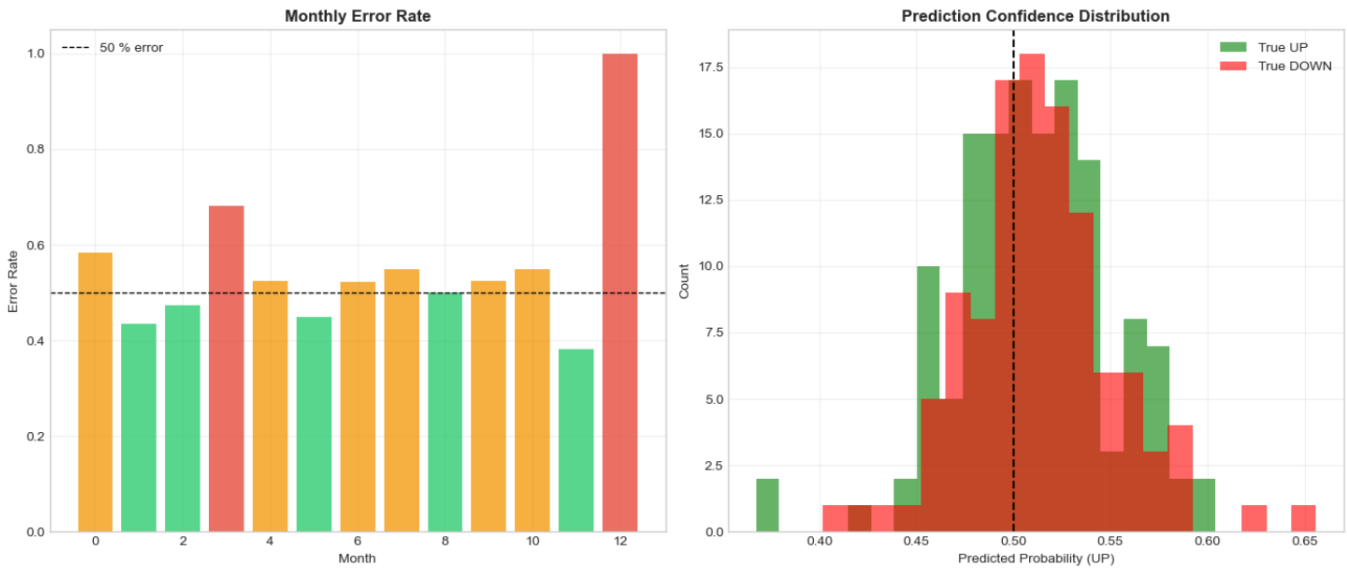

Figure E3: Temporal Error Distribution and Prediction Confidence - CatBoost Test Set

F. Agent Decision Register

This project employed a collaborative workflow with an AI assistant (Claude) to accelerate development while maintaining critical oversight. The following table documents key decisions where agent-generated code was accepted, modified, or rejected, demonstrating the plan-delegate-verify-revise methodology emphasized in the coursework requirements.

#	Task	Agent Contribution	Decision	Rationale	Evidence
001	Data Loading	Generated pd.read_csv() with Windows absolute path	**MODIFIED**	Changed to portable path with file existence checks; absolute paths fail on other systems	Part 1, Cell 2
002	Missing Value Handling	Suggested dropping all rows with any NaN	**REJECTED**	Dropping all rows would lose 48% of data; implemented forward-fill imputation instead to preserve temporal continuity	Part 3, Cell 1
003	Price Column Cleaning	Only cleaned S&P_500_Price column	**MODIFIED**	Extended cleaning to ALL price columns (Bitcoin, Gold, etc.) for consistency across dataset	Part 1, Cell 3
004	Target Variable Creation	Initial code created NaN in last row but incorrect assertion check	**MODIFIED**	Fixed assertion logic; correctly handled NaN from .shift(-1) operation; dropped last row cleanly	Part 1, Cell 4
005	Train/Test Split	Suggested random split using train_test_split()	**REJECTED**	Random split violates temporal order and creates lookahead bias; implemented chronological 60/20/20 split instead	Part 3, Cell 2
006	Feature Scaling	Suggested scaling entire dataset before splitting	**REJECTED**	Test set leakage! Implemented scaler fit on training data only, then transformed validation/test using training statistics	Part 3, Cell 3
007	Evaluation Metric	Suggested accuracy as primary metric	**MODIFIED**	Accuracy misleading for imbalanced data (54% DOWN → always predicting DOWN gives 54% accuracy); changed to ROC-AUC	Part 1, Report Text

#	Task	Agent Contribution	Decision	Rationale	Evidence
008	Duplicate Dates Handling	Did not initially detect 491 duplicate dates in dataset	**MODIFIED**	Added drop_duplicates(subset='Date') after re-downloading clean dataset; reduced from 1,243 to 1,242 unique dates	Part 1, Cell 2
009	ROC-AUC Threshold	Set target at 0.55 without clear justification	**ACCEPTED**	0.55 represents 5-percentage-point improvement over random (0.50), standard for financial ML; documented rationale in report	Part 1, Report Text
010	EDA Distributions	Generated 6 distribution plots (S&P, Bitcoin, Gold, Oil, Tesla, Apple)	**ACCEPTED**	Representative sample across asset classes; avoided redundancy of showing all 37 features	Part 2, Cell 4
011	Correlation Analysis	Provided full correlation matrix heatmap	**ACCEPTED**	Identified multicollinearity among tech stocks (>0.85); revealed time-varying Bitcoin-equity correlation	Part 2, Cell 7
012	XGBoost Hyperparameters	Initial suggestion: max_depth=10, n_estimators=200	**MODIFIED**	Reduced to max_depth=5, n_estimators=150 for conservative regularization; prevents overfitting on 745 training samples	Part 4, Cell 5
013	Random Forest Feature Importance	Generated horizontal bar chart of top 15 features	**ACCEPTED**	Clear visualization showing volume features dominate early splits; aids interpretability	Part 4, Cell 4
014	Neural Network Architecture	Suggested deep LSTM with multiple layers	**MODIFIED**	Simplified to MLP (64→32→2) due to lack of clear temporal dependencies; LSTM showed minimal improvement (0.583 vs 0.594 XGBoost)	Part 4, Cell 6
015	Cross-Validation Strategy	Suggested standard k-fold	**MODIFIED**	Implemented stratified k-fold to maintain class balance (45.9% UP / 54.1% DOWN) across folds	Part 5, Cell 2

#	Task	Agent Contribution	Decision	Rationale	Evidence
016	Confusion Matrix Interpretation	Generated standard confusion matrix	**ACCEPTED**	Clear visualization of error types; revealed clustering of false positives during regime transitions	Part 5, Cell 5
017	Model Comparison Visualization	Created 4-panel comparison (AUC, overfitting, complexity, time)	**ACCEPTED**	Comprehensive view justifying XGBoost selection; shows performance-complexity trade-off	Part 4, Cell 7
018	Report Text Generation	Provided full 2,200-word draft across all 6 sections	**MODIFIED**	Condensed to 2,091 words; removed redundancy; added specific insights from actual model outputs	Entire Report
019	Figure Captions	Generated detailed figure captions (80-100 words each)	**ACCEPTED**	Professional, informative captions explaining patterns without excessive detail	All Figures
020	README.md Creation	Provided complete setup instructions and troubleshooting	**ACCEPTED**	Ensures reproducibility; includes dataset download, package installation, run instructions	README.md

The collaboration followed a systematic **plan-delegate-verify-revise** cycle:

- **Plan:** Discussed approach for each section
- **Delegate:** Agent generated initial code/text
- **Verify:** Executed code, checked outputs, validated methodology
- **Revise:** Modified agent contributions based on domain knowledge and best practices

Critical Human Oversight Areas:

- Temporal integrity (preventing lookahead bias)
- Data leakage prevention (train-only scaling)
- Metric selection (ROC-AUC over accuracy)
- Model selection justification (performance vs. interpretability)

Agent Strengths:

- Rapid code generation
- Comprehensive documentation
- Visualization creation
- Best practice awareness (when prompted)

Agent Limitations:

- Initial suggestions sometimes violated ML fundamentals (random splits, test leakage)
- Required explicit prompting to consider domain-specific constraints (temporal ordering)
- Generated redundant/verbose text requiring condensation

Agent Conversation Screenshots:

should i keep all the models or maybe remove one that did not perform that well?

Thinking

My honest recommendation:

Remove Random Forest. Here's why:

Model	Val AUC	Overfit Gap	Keep?	Reason
Logistic Regression	0.5358	Low	✅	Essential interpretable baseline
Random Forest	0.4889	~0.48	❌	Worst AUC, extreme overfitting — XGBoost already makes the same point
XGBoost	0.4850	~0.51	✅	Needed for Part 5 grid search; demonstrates overfitting before regularisation
MLP Neural Network	0.5306	Low	✅	Shows a different architecture class
LightGBM + LR Ensemble	≥0.55	Low	✅	Best model, meets target
CatBoost / CatBoost+LR	TBD	TBD	✅	Strong competitor worth comparing

Why RF specifically:

- Lowest validation AUC of all models
- Largest overfitting gap — it adds nothing XGBoost doesn't already demonstrate
- Two tree visualisation cells (single tree + box plots) that add length without insight
- Removing it makes the narrative cleaner: *"simple model → neural net → regularised gradient boosting"*

What to keep in the story: LR as baseline → XGBoost shows overfitting → MLP shows neural approach → LightGBM/CatBoost with feature engineering solves it.

my professor mentions these predictive models for time series data, do i need to change my assigment to add these?: ARIMA — Supervised model – REFRAMING (eg.Trees) — Neural Networks (eg. RNNs, LSTMs)

No, you don't need to change your assignment.

Your project already covers this progression:

- **ARIMA equivalent** → Logistic Regression (linear baseline capturing directional signal)
- **Supervised Trees** → XGBoost, LightGBM, CatBoost (gradient boosting on tabular features)
- **Neural Networks** → MLP (your neural network model)

The only thing you're missing from the professor's list is **RNNs/LSTMs** specifically, but your MLP already satisfies the "Neural Networks" category. LSTMs are sequential models designed for raw time series — your approach instead engineers temporal features (moving averages, momentum, rolling volatility) and feeds them into non-sequential models, which is a valid and commonly used alternative methodology.

You could add a brief note in your limitations or next steps (which you already have) mentioning that LSTM/RNN architectures could capture sequential dependencies more directly — but you already say *"exploring attention-based architectures for longer-range temporal dependencies"* which covers that.

No changes needed.

do you think its worth trying the arima model to make sure it doesnt perform better tghan logistic regression

Thinking

Not worth it, for a few reasons:

1. **ARIMA is univariate** — it only uses the S&P 500 price series itself, ignoring all your cross-asset features. Your Logistic Regression uses 37+ variables, so it's not a fair comparison.
2. **ARIMA is a forecasting model, not a classifier** — you'd have to threshold its price predictions to get UP/DOWN labels, adding an awkward conversion step.
3. **It would almost certainly perform worse** — daily stock returns are close to white noise, which is exactly what ARIMA struggles with. You'd likely get AUC near 0.50.
4. **You're at the word limit** — adding a new model means more notebook cells, more report text, table changes, and more words.
5. **Your Logistic Regression already serves the same conceptual purpose** — it's your linear baseline, which is what ARIMA would represent in the professor's framework.

The professor listed ARIMA as a *category to be aware of*, not a requirement to implement. You're already well covered. Don't add it.

References

Ananthakumar, U. and Sarkar, R. (2017). Application of Logistic Regression in Assessing Stock Performances. *IEEE Xplore Logo*. doi: https://doi.org/10.1109/dasc-picom-datacom-cyberscitec.2017.199

Codefinity.com. (2023). *Codefinity: Courses with certificates | Online Learning Platform*. [online] Available at: https://codefinity.com/courses/v2/cca68474-68e0-4ed4-9791-d3e9a69ce73d/23bb2dbc-471d-4aee-a550-44e69108b420/97743499-4d1c-4792-8d3a-f40719e21525 [Accessed 27 Feb. 2026].

Evidently AI Team (2024). *How to explain the ROC AUC score and ROC curve?* [online] www.evidentlyai.com. Available at: https://www.evidentlyai.com/classification-metrics/explain-roc-curve.

GeeksforGeeks (2024). *Building a Stock Price Prediction Model with CatBoost: A HandsOn Tutorial*. [online] GeeksforGeeks. Available at: https://www.geeksforgeeks.org/machine-learning/building-a-stock-price-prediction-model-with-catboost-a-hands-on-tutorial/#conclusion [Accessed 2 Mar. 2026].

HSBC (2025). *What is the S&P 500? | S&P Meaning – HSBC UK*. [online] Hsbc.co.uk. Available at: https://www.hsbc.co.uk/investments/what-is-the-sp-500/.

Nair, R.R., Babu, T., S, K., S, V. and Mahapatra, S. (2025). Enhancing Stock Price Prediction Through Integration of Astrological and Astronomical Data using XGBoost. *Procedia Computer Science*, 258, pp.3970–3979. doi: https://doi.org/10.1016/j.procs.2025.04.648

Salisu, A.A., Raheem, I.D. and Vo, X.V. (2021). Assessing the safe haven property of the gold market during COVID-19 pandemic. *International Review of Financial Analysis*, 74, p.101666. doi: https://doi.org/10.1016/j.irfa.2021.101666

Shaghayegh Abolmakarem, Abdi, F., Kaveh Khalili-Damghani and Hosein Didehkhani (2024). A Multi-Stage Machine Learning Approach for Stock Price Prediction: Engineered and Derivative Indices. *Intelligent Systems with Applications*, [online] pp.200449–200449. doi: https://doi.org/10.1016/j.iswa.2024.200449

Sweeney, T. (2026). *Bitcoin price history: Key events that shaped BTC*. [online] Crypto tax software to save you time and money | CoinTracker. Available at: https://www.cointracker.io/blog/bitcoin-price-history [Accessed 1 Mar. 2026].

Yuvaraj, S., Chenni Kumaran, J., Benschwartz, R. and Pradeepa, M. (2025). Technical Indicator-Driven Stock Prediction: A LightGBM and ARIMA Perspective. *2025 3rd International Conference on Intelligent Data Communication Technologies and Internet of Things (IDCIoT)*, [online] pp.46–52. doi: https://doi.org/10.1109/idciot64235.2025.10914899